

Full Name	Student ID
Michael Ekubay Negassi	002238528
Lu Yang	002238352

Statistical Learning Summer Project

1.

Prerequisites

Install these packages for and libraries for Question 1.

```
install.packages("questionr")
library(questionr)
```

Load **lowbwt.txt** file from saved Directory

```
temp<-read.table("C:/Challenge2/lowbwt.txt",header = TRUE)
```

#a)

Write the equation and interpret β_1 , then estimated coefficient of Apgar score.

```
apgar5<-temp$apgar5
germ.hem<-temp$germ.hem
apgar5.glm = glm(formula=germ.hem ~ apgar5, data=temp, family=binomial
(link="logit"))
summary(apgar5.glm)
```

```
Call:
glm(formula = germ.hem ~ apgar5, family = binomial(link = "logit"
),
    data = temp)
```

Deviance Residuals:

```
      Min       1Q   Median       3Q      Max
-1.0514  -0.5528  -0.4645  -0.3877   2.1891
```

Coefficients:

```
              Estimate Std. Error z value Pr(>|z|)
(Intercept)  -0.3037     0.6191  -0.491   0.6237
apgar5        -0.2496     0.1044  -2.392   0.0168 *
```

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 84.542 on 99 degrees of freedom

Residual deviance: 78.927 on 98 degrees of freedom

AIC: 82.927

Number of Fisher Scoring iterations: 4

Equation : $\ln(p/1-p) = -0.3037 - 0.2496x$

For each one unit of increase in the apgar score, the log odds of experiencing a germinal matrix hemorrhage decrease by 0.2496

#b)

confint(apgar5.glm)

		2.5 %	97.5 %
(Intercept)	-1.5832887		0.89613184
apgar5	-0.4599101		-0.04420464

exp(coef(apgar5.glm))

	(Intercept)	apgar5
	0.7380593	0.7791066

exp(cbind(OR=coef(apgar5.glm),confint(apgar5.glm)))

	OR	2.5 %	97.5 %
(Intercept)	0.7380593	0.2052988	2.4501073
apgar5	0.7791066	0.6313404	0.9567581

When a logistic regression is calculated, the regression coefficient (β_1) is the estimated increase in the log odds of the *outcome per unit increase* in the value of the *exposure*. In other words, the exponential function of the regression coefficient ($e\beta_1$) is the odds ratio associated with a one-unit increase in the exposure. Odds ratio could also be obtained with $\exp(\text{coef}(x))$ and confidence intervals with $\exp(\text{confint}(x))$.

#c)

H0: $\hat{\tau}^2_1=0$ Logit graham apgar5

Iteration 0: log likelihood = -42.270909
 Iteration 1: log likelihood = -39.727053
 Iteration 2: log likelihood = -39.463638
 Iteration 3: log likelihood = -39.463411

Logistic regression Log likelihood = -39.463411

#d)

new.apgar<-data.frame(apgar5=3)***prob<-predict(apgar5.glm,new.apgar,type = "response")******prob***

0.2587351

new.apgar2<-data.frame(apgar5=7)***prob2<-predict(apgar5.glm,new.apgar2,type = "response")******prob2***

0.1139531

2.

Prerequisites

#Install these packages for and libraries for Question 2.

install.packages("ISLR")***library(ISLR)******install.packages("corrplot")***

Load Auto CSV file to variable AUTO and

Clearing data with "?" values from source horsepower column

Auto<-read.csv("C:/Challenge2/Auto.csv",header = TRUE,sep = ',')***Auto\$horsepower <- as.numeric(as.character(Auto\$horsepower))******Auto <- na.omit(Auto)***

#a)

mpg01 <- rep(0, length(Auto\$mpg))***mpg01[Auto\$mpg > median(Auto\$mpg)] <- 1******Auto <- data.frame(Auto, mpg01)******summary(Auto)***

```

      mpg      cylinders      displacement      horsepower      weight      acceleration      year      origin      name
Min.   : 9.00   Min.   :3.000   Min.   : 68.0   Min.   : 46.0   Min.   :1613   Min.   : 8.00   Min.   :70.00   Min.   :1.000   amc matador      : 5
1st Qu.:17.00   1st Qu.:4.000   1st Qu.:105.0   1st Qu.: 75.0   1st Qu.:2225   1st Qu.:13.78   1st Qu.:73.00   1st Qu.:1.000   ford pinto      : 5
Median :22.75   Median :4.000   Median :151.0   Median : 93.5   Median :2804   Median :15.50   Median :76.00   Median :1.000   toyota corolla   : 5
Mean   :23.45   Mean   :5.472   Mean   :194.4   Mean   :104.5   Mean   :2978   Mean   :15.54   Mean   :75.98   Mean   :1.577   amc gremlin     : 4
3rd Qu.:29.00   3rd Qu.:8.000   3rd Qu.:275.8   3rd Qu.:126.0   3rd Qu.:3615   3rd Qu.:17.02   3rd Qu.:79.00   3rd Qu.:2.000   amc hornet      : 4
Max.   :46.60   Max.   :8.000   Max.   :455.0   Max.   :230.0   Max.   :5140   Max.   :24.80   Max.   :82.00   Max.   :3.000   chevrolet chevette: 4
                                           (other)          :365

      mpg01
Min.   :0.0
1st Qu.:0.0
Median :0.5
Mean   :0.5
3rd Qu.:1.0
Max.   :1.0

```

#b) This depends on the mpg01 column created above

```
cor(Auto[, -9])
```


```

      mpg      cylinders      displacement      horsepower      weight      acceleration      year      origin      mpg01
mpg      1.0000000   -0.7776175   -0.8051269   -0.7784268   -0.8322442   0.4233285   0.5805410   0.5652088   0.8369392
cylinders -0.7776175   1.0000000   0.9508233   0.8429834   0.8975273   -0.5046834   -0.3456474   -0.5689316   -0.7591939
displacement -0.8051269   0.9508233   1.0000000   0.8972570   0.9329944   -0.5438005   -0.3698552   -0.6145351   -0.7534766
horsepower -0.7784268   0.8429834   0.8972570   1.0000000   0.8645377   -0.6891955   -0.4163615   -0.4551715   -0.6670526
weight      -0.8322442   0.8975273   0.9329944   0.8645377   1.0000000   -0.4168392   -0.3091199   -0.5850054   -0.7577566
acceleration 0.4233285   -0.5046834   -0.5438005   -0.6891955   -0.4168392   1.0000000   0.2903161   0.2127458   0.3468215
year         0.5805410   -0.3456474   -0.3698552   -0.4163615   -0.3091199   0.2903161   1.0000000   0.1815277   0.4299042
origin       0.5652088   -0.5689316   -0.6145351   -0.4551715   -0.5850054   0.2127458   0.1815277   1.0000000   0.5136984
mpg01       0.8369392   -0.7591939   -0.7534766   -0.6670526   -0.7577566   0.3468215   0.4299042   0.5136984   1.0000000

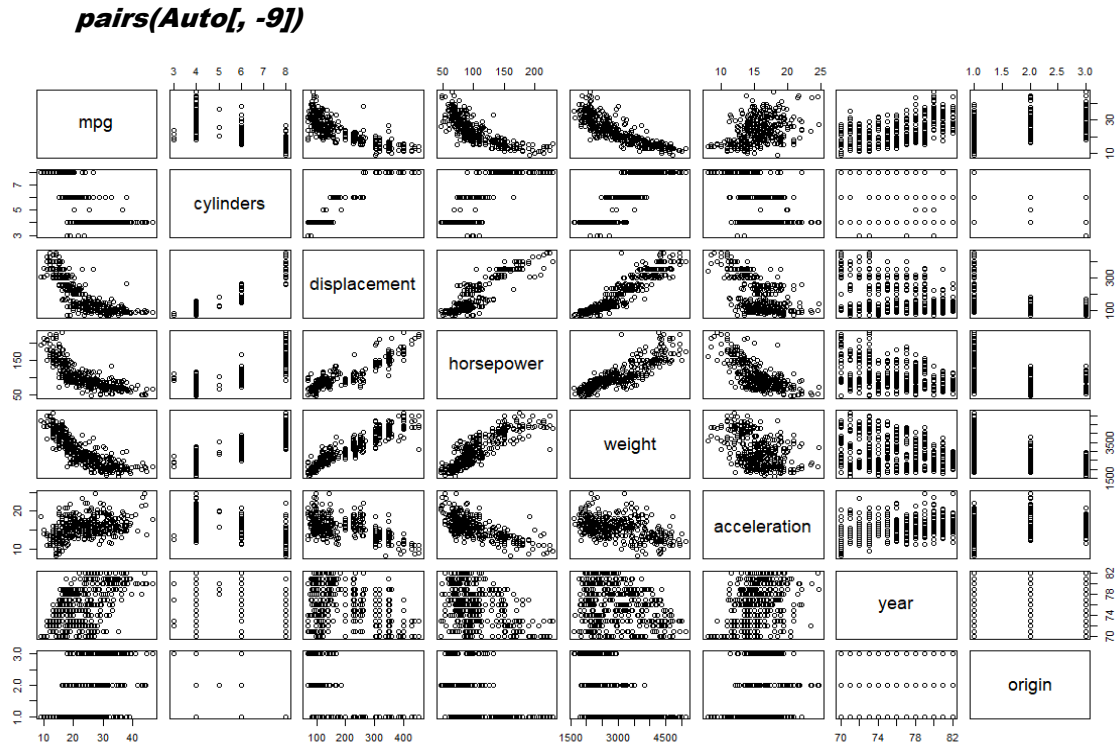
```

```
library(corrplot)
```

```
corrplot::corrplot.mixed(cor(Auto[, -9]), upper="circle")
```



Scatterplot matrix

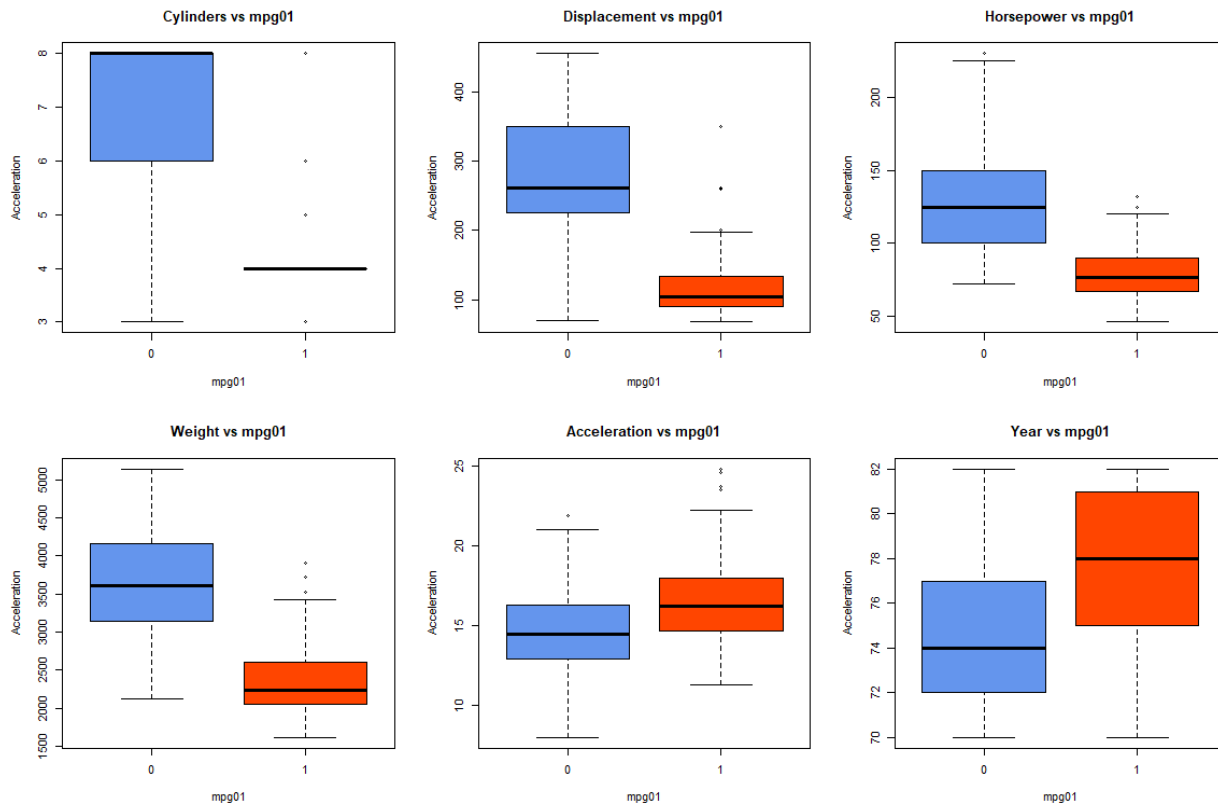


BoxPlots

```

par(mfrow=c(2,3))
boxplot(cylinders ~ mpg01, data = Auto, main = "Cylinders vs mpg01" ,
xlab = "mpg01", ylab = "Acceleration",col = c("cornflowerblue", "orangered"))
boxplot(displacement ~ mpg01, data = Auto, main = "Displacement vs mpg01" ,
xlab = "mpg01", ylab = "Acceleration",col = c("cornflowerblue", "orangered"))
boxplot(horsepower ~ mpg01, data = Auto, main = "Horsepower vs mpg01" ,
xlab = "mpg01", ylab = "Acceleration",col = c("cornflowerblue", "orangered"))
boxplot(weight ~ mpg01, data = Auto, main = "Weight vs mpg01" , xlab = "mpg01",
ylab = "Acceleration",col = c("cornflowerblue", "orangered"))
boxplot(acceleration ~ mpg01, data = Auto, main = "Acceleration vs mpg01" ,
xlab = "mpg01", ylab = "Acceleration",col = c("cornflowerblue", "orangered"))
boxplot(year ~ mpg01, data = Auto, main = "Year vs mpg01" , xlab = "mpg01",
ylab = "Acceleration",col = c("cornflowerblue", "orangered"))

```



Based on the above plots, mpg01 has a negative association with Weight, Cylinders, Displacement, and Horsepower.

#c) Split the data into a training set and a test set.

```
set.seed(123)
train <- sample(1:dim(Auto)[1], dim(Auto)[1]*.7, rep=FALSE)
test <- -train
Training_dataSet<- Auto[train, ]
Testing_dataSet= Auto[test, ]
mpg01.test <- mpg01[test]
```

#d) LDA on the training data to predict mpg01 using the variables that seemed most associated with mpg01

```
library(MASS)
lda_model <- lda(mpg01 ~ cylinders + weight + displacement + horsepower, dat
a = Training_dataSet)
lda_model
```

```
[1] 0.1186441
```

Using all 4 predictors (cylinders, weight, displacement, and horsepower) with a Linear Discriminant Analysis the test error rate is 11.86441%.

#f)

```
glm_model <- glm(mpg01 ~ cylinders + weight + displacement + horsepower, data = Training_dataSet, family = binomial)
summary(glm_model)
```

Call:

```
glm(formula = mpg01 ~ cylinders + weight + displacement + horsepower,
     family = binomial, data = Training_dataSet)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.4870	-0.1763	0.1231	0.3633	3.2024

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	11.6661375	1.9885822	5.867	4.45e-09	***
cylinders	0.0365940	0.3817545	0.096	0.92363	
weight	-0.0023706	0.0008319	-2.850	0.00438	**
displacement	-0.0106969	0.0096956	-1.103	0.26991	
horsepower	-0.0327332	0.0154688	-2.116	0.03434	*

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 379.48 on 273 degrees of freedom
 Residual deviance: 147.50 on 269 degrees of freedom
 AIC: 157.5

Number of Fisher Scoring iterations: 7

```
probs <- predict(glm_model, Testing_dataSet, type = "response")
pred.glm <- rep(0, length(probs))
pred.glm[probs > 0.5] <- 1
table(pred.glm, mpg01.test)
```

	mpg01.test	
pred.glm	0	1
0	54	3
1	10	51

```
mean(pred.glm != mpg01.test)
```

```
[1] 0.1101695
```

Using all 4 predictors (cylinders, weight, displacement, and horsepower) with a Logistic Regression Analysis the test error rate is 11.01695 %.

#g)

str(Auto)

```
'data.frame': 392 obs. of 10 variables:
 $ mpg      : num  18 15 18 16 17 15 14 14 14 15 ...
 $ cylinders : int   8  8  8  8  8  8  8  8  8  8 ...
 $ displacement: num  307 350 318 304 302 429 454 440 455 390 ...
 $ horsepower  : num  130 165 150 150 140 198 220 215 225 190 ...
 $ weight      : int  3504 3693 3436 3433 3449 4341 4354 4312 4425 3850 ...
 $ acceleration: num   12 11.5 11 12 10.5 10 9 8.5 10 8.5 ...
 $ year        : int   70  70  70  70  70  70  70  70  70  70 ...
 $ origin      : int    1  1  1  1  1  1  1  1  1  1 ...
 $ name        : Factor w/ 304 levels "amc ambassador brougham",...: 49 36 231
14 161 141 54 223 241 2 ...
 $ mpg01      : num   0  0  0  0  0  0  0  0  0  0 ...
```

data = scale(Auto[, -c(9, 10)])***set.seed(1234)******train <- sample(1:dim(Auto)[1], 392*.7, rep=FALSE)******test <- -train******Training_dataSet = data[train, c("cylinders", "horsepower", "weight", "acceleration")]******Testing_dataSet = data[test, c("cylinders", "horsepower", "weight", "acceleration")]******## KNN take the training response variable seperately******train.mpg01 = Auto\$mpg01[train]******## we also need the have the testing_y seperately for assesing the model later on******test.mpg01 = Auto\$mpg01[test]******library(class)******set.seed(1234)******knn_pred_y = knn(Training_dataSet, Testing_dataSet, train.mpg01, k = 1)******table(knn_pred_y, test.mpg01)***

```
      test.mpg01
knn_pred_y 0  1
0    50  4
1     6 58
```

mean(knn_pred_y != test.mpg01)

[1] 0.08474576

Test error rate is 8.474576 % for K=1.

```
knn_pred_y = NULL  
error_rate = NULL  
for(i in 1:dim(Testing_dataSet)[1]){  
  set.seed(1234)  
  knn_pred_y = knn(Training_dataSet,Testing_dataSet,train.mpg01,k=i)  
  error_rate[i] = mean(test.mpg01 != knn_pred_y)  
}  
### find the minimum error rate  
min_error_rate = min(error_rate)  
print(min_error_rate)
```

```
[1] 0.06779661
```

K with lowest values

```
K = which(error_rate == min_error_rate)  
print(K)
```

```
[1] 4
```

When we train a KNN model with k=4 then we get the lowest misclassification error rate of 6.779661%.

```
library(ggplot2)  
qplot(1:dim(Testing_dataSet)[1], error_rate, xlab = "K", ylab = "Error Rate",  
geom=c("point", "line"))
```

